

CS 4602

Introduction to Machine Learning

Decision Trees
(Bagging and Boosting)

Instructor: Po-Chih Kuo

Roadmap

- Introduction and Basic Concepts
- Regression (Error-Based Learning)
- Bayesian Classifiers (Probability-Based Learning)
- Decision Trees (Information-Based Learning)
- KNN (Similarity-Based Learning)
- Linear Classifier
- Neural Networks
- Deep learning
- Convolutional Neural Networks
- RNN/Transformer
- Reinforcement Learning
- Model Selection and Evaluation
- Clustering
- Data Exploration & Dimensionality reduction

Make a Decision




Problem: decide whether to wait for a table at a restaurant based on the following **attributes**:

1. **Alternate**: is there an alternative restaurant nearby?
2. **Bar**: is there a comfortable bar area to wait in?
3. **Fri/Sat**: is today Friday or Saturday?
4. **Hungry**: are we hungry?
5. **Patrons**: number of people in the restaurant (None, Some, Full)
6. **Price**: price range (\$, \$\$, \$\$\$)
7. **Raining**: is it raining outside?
8. **Reservation**: have we made a reservation?
9. **Type**: kind of restaurant (French, Italian, Thai, Burger)
10. **WaitEstimate**: estimated waiting time (0-10, 10-30, 30-60, >60)

Attribute-based representations

- Examples described by **attribute values** (Boolean, discrete, continuous)



Example	Attributes										Target
	Alt.	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est.	Wait
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30-60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0-10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10-30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0-10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0-10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30-60	T

d

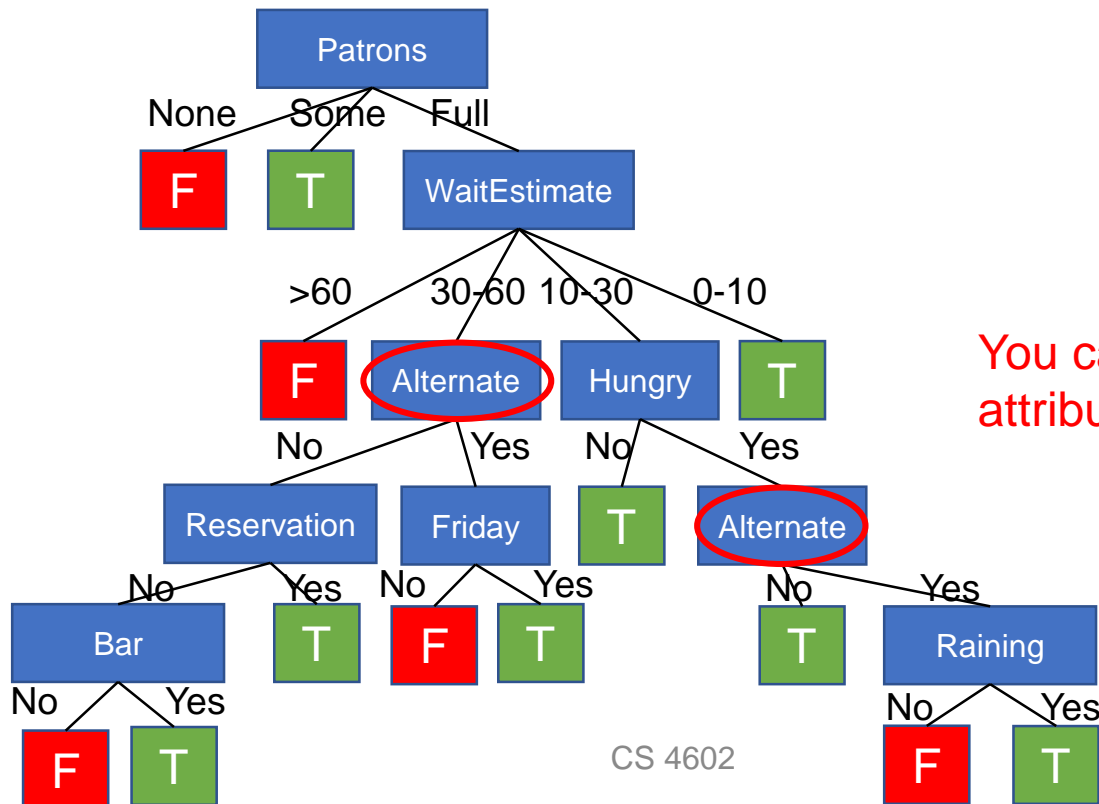
y

- Classification** of examples is **positive** (T: wait) or **negative** (F: Not to wait)
- A number N of instances, each with attributes $(x_1, x_2, x_3, \dots, x_d)$ and target value y .

Decision trees



- One possible representation for **hypotheses**
- We imagine someone taking a sequence of decisions.
- E.g., here is the **Optimal** tree for deciding whether to wait:

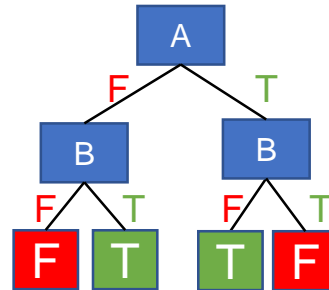


You can use the same attribute more than once.

Expressiveness

- Decision trees can express any function of the input attributes.
- E.g., for Boolean functions, truth table row \rightarrow path to leaf:

A	B	A xor B
F	F	F
F	T	T
T	F	T
T	T	F



- Trivially, there is a consistent decision tree for any training set with one path to leaf for each example, but **it probably won't generalize to new examples.**
OVERFITTING!
- Prefer to find more **compact** decision trees: we don't want to memorize the data. We want to find **structure** in the data!



Hypothesis spaces

How many distinct decision trees with d Boolean attributes?

= number of Boolean functions

= number of distinct truth tables with 2^d rows = 2^{2^d}

X_1	X_2	Y
0	0	?
0	1	?
1	0	?
1	1	?



$n=2$: $2^2 = 4$ rows. For each row we can choose T or F: 2^4 functions.

- E.g., with 6 Boolean attributes, how many trees?


There are 18,446,744,073,709,551,616 trees !

Decision tree learning

- If there are so many possible trees, can we actually search this space? (solution: greedy search).
- **Aim**: find a small tree consistent with the training examples
- **Idea**: Recursively choose "most significant" attribute as root of (sub-) tree.

Choosing a significant attribute

- Idea: a good attribute splits the examples into subsets that are (ideally) "all positive" or "all negative"



To wait or not to wait is still at 50%!

- Patrons or type?*


Information theory

- **Entropy** measures the amount of uncertainty in a **probability** distribution:
 - Consider tossing a **biased** coin.
 - If you toss the coin VERY often, the frequency of heads is p and the frequency of tails is $1-p$. (fair coin $p=0.5$).
 - The uncertainty in any actual outcome is given by the entropy.
 - The uncertainty is zero if $p=0$ or 1 and maximal if we have $p=0.5$.



$$\text{Entropy}(X) = -p \log_2 p - (1 - p) \log_2 (1 - p)$$

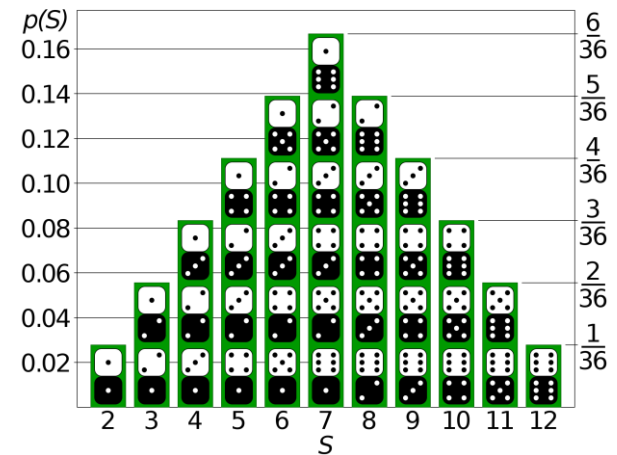
binary entropy function



Information theory

- If there are more than two states $s=1,2,..n$ we have (e.g. a die):

$$\begin{aligned}\text{Entropy}(X) = & \\ & -p(s = 1) \log_2[p(s = 1)] \\ & -p(s = 2) \log_2[p(s = 2)] \\ & \dots \\ & -p(s = n) \log_2[p(s = n)]\end{aligned}$$



$$\sum_{s=1}^n p(s) = 1$$

Information theory

- Imagine we have training data: p positive examples and n negative examples.
- Our best estimate of true or false is given by:

$$P(\text{positive}) \approx \frac{p}{p+n}$$
$$P(\text{negative}) \approx \frac{n}{p+n}$$

- Hence the entropy is given by:

$$\text{Entropy}\left(\frac{p}{p+n}, \frac{n}{p+n}\right) \approx -\frac{p}{p+n} \log \frac{p}{p+n} - \frac{n}{p+n} \log \frac{n}{p+n}$$

What about **Cross-entropy**?

- **Entropy** is a measure of the uncertainty in **one** probability distribution.
- **Cross-entropy** is a measure of the difference between **two** probability distributions.



We'll talk about this more in the neural network section

Information gain

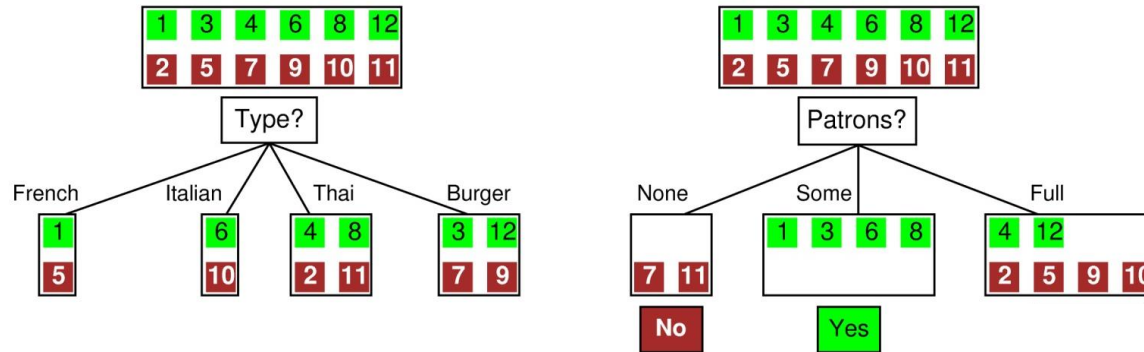
- How much information do we gain if we disclose the value of some attribute?
- Answer:

Decrease of uncertainty

=uncertainty before - uncertainty after

=entropy before - entropy after

Example



Before: Entropy = $-\frac{1}{2} \log(1/2) - \frac{1}{2} \log(1/2) = \log(2) = 1$ bit:
There is “1 bit of information to be discovered”.

After: for **Type**: If we go into branch “French” we have 1 bit, similarly for the others.

French: 1bit
Italian: 1 bit
Thai: 1 bit
Burger: 1bit

} On average: 1 bit ! We gained nothing!

After: for **Patrons**: In branch “None” and “Some” entropy = 0
In branch “Full” entropy = $-\frac{1}{3} \log(1/3) - \frac{2}{3} \log(2/3) = 0.918....$

So Patrons gains more information!

Information gain


- How do we combine branches:

1/6 of the time we enter “None”, so we weight “None” with 1/6. Similarly: “Some” has weight: 1/3 and “Full” has weight 1/2.

$$Entropy(A) = \sum_{i=1}^n \frac{p_i + n_i}{p + n} Entropy\left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}\right)$$

weight for each branch

entropy for each branch.



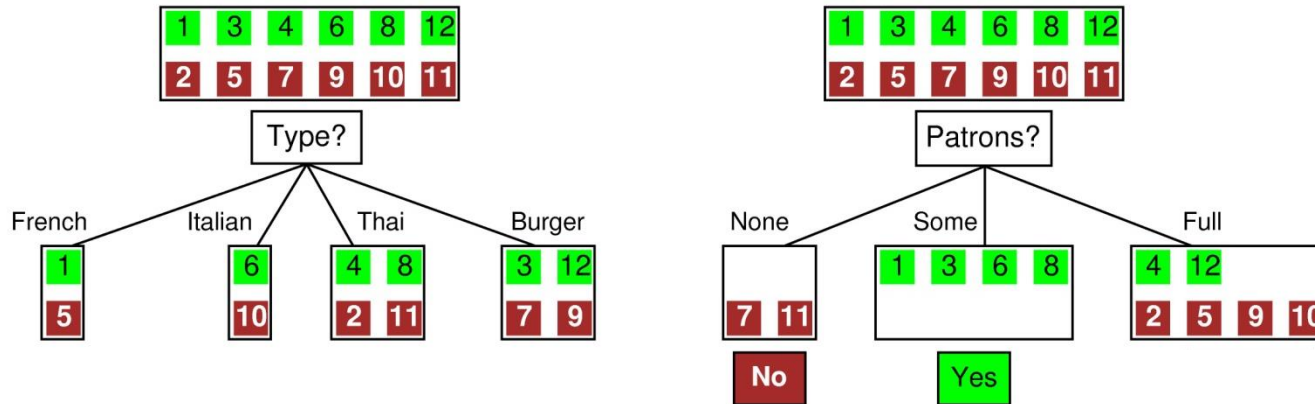
Information gain

- **Information Gain (IG)** or reduction in entropy from the attribute test:

$$IG(A) = Entropy\ before - Entropy\ after$$

- Choose the attribute with the **largest** IG

Information gain



For the training set, $p = n = 6$, $E(6/12, 6/12) = 1$ bit

$$IG(Patrons) = 1 - \left[\frac{2}{12} E(0,1) + \frac{4}{12} E(1,0) + \frac{6}{12} E\left(\frac{2}{6}, \frac{4}{6}\right) \right] = 0.541 \text{ bits}$$

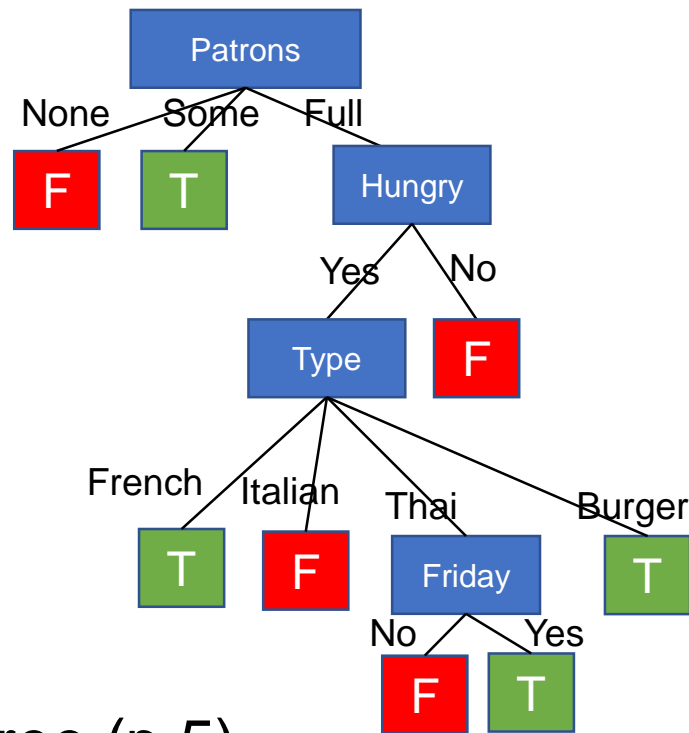
$$IG(Type) = 1 - \left[\frac{2}{12} E\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{2}{12} E\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{4}{12} E\left(\frac{2}{4}, \frac{2}{4}\right) + \frac{4}{12} E\left(\frac{2}{4}, \frac{2}{4}\right) \right] = 0 \text{ bits}$$

Patrons has the highest IG of all attributes and so is chosen by the Decision Tree Learning algorithm as the root

Is it sensible to have the same attribute on a single branch of the tree (why)?


Example contd.

- Decision tree learned from the 12 examples:



- Simpler than **TRUE** tree (p.5)
 - a more complex hypothesis isn't justified by small amount of data

Gain-Ratio



- If 1 attribute splits in many more classes than another, it has an (unfair) advantage if we use information gain.
- The gain-ratio is designed to compensate for this problem,

$$\text{Gain Ratio} = \frac{IG}{-\sum_{i=1}^n \frac{p_i + n_i}{p + n} \log \frac{p_i + n_i}{p + n}}$$

- if we have n uniformly populated classes the denominator is $\log_2(n)$ which penalized relative to 1 for 2 uniformly populated classes.

Gain-Ratio (Example)

		Wait		
		T	F	Total
Patrons	Full	0	2	2
	Some	4	0	4
	None	2	4	6

$$IG(Patrons) = 0.541$$

$$\begin{aligned} \text{Gain Ratio} &= \frac{IG}{-(\frac{2}{12}\log_2(\frac{2}{12}) + \frac{4}{12}\log_2(\frac{4}{12}) + \frac{6}{12}\log_2(\frac{6}{12}))} \\ &= \frac{0.541}{1.459} = 0.371 \end{aligned}$$



What to Do if...

- In some leaf there are no examples:

Choose True or False according to the number of positive/negative examples at your parent.


- There are no attributes left

Two or more examples have the same attributes but different label: we have an error/noise. Stop and use majority vote.



Continuous variables

- If variables are continuous we can
 - bin them



- learn a simple classifier on a single dimension (e.g. logistic regression classifier).

A practical way for continuous variables

Ex.	Est	Wait
X ₁	8	T
X ₂	40	F
X ₃	8	T
X ₄	12	T
X ₅	70	F
X ₆	3	T
X ₇	7	F
X ₈	6	T
X ₉	80	F
X ₁₀	20	F
X ₁₁	8	F
X ₁₂	40	T

sort
➔

Ex.	Est	Wait
X ₆	3	T
X ₈	6	T
X ₇	7	F
X ₁	8	T
X ₃	8	T
X ₁₁	8	F
X ₄	12	T
X ₁₀	20	F
X ₂	40	F
X ₁₂	40	T
X ₅	70	F
X ₉	80	F

Middle point (what else?)


4.5 ➔ IG
6.5 ➔ IG
7.5 ➔ IG


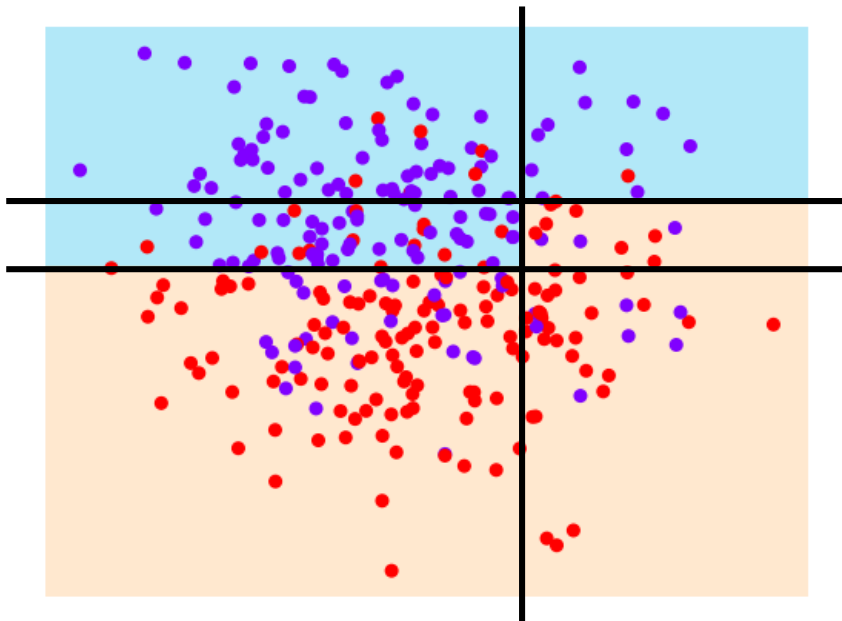


10 ➔ IG
16 ➔ IG
30 ➔ IG

55 ➔ IG
75 ➔ IG

if there are **N** possible values, we would have at most **N-1** possible splits.

➔ Max(IGs) is the IG for this attribute (EST.)





When to Stop ?


- If we keep going until perfect classification we might over-fit.
- Heuristics:
 - Stop when Info-Gain (Gain-Ratio) is smaller than threshold
 - Stop when there are M examples in each leaf node
- Penalize complex trees by minimizing with “complexity” = # nodes.
Note: if tree grows, complexity grows but entropy shrinks.

$$\alpha \times \text{complexity} + \sum_{\text{all leaves}} \text{entropy}(\text{leaf})$$

- Compute many full grown trees on subsets of data and test them on hold-out data. Pick the best or average their prediction.
- Do a statistical test: is the increase in information significant.

How to improve Decision Trees?

“Unity is strength!”



Ensemble methods!

Ensemble methods: bagging, boosting

Aims

- **Bagging (bootstrap aggregating)** algorithms aim to reduce the complexity of models that **overfit** the training data.
- **Boosting** is an approach to increase the complexity of models that **underfit** the training data.

How do they work?

- **Bagging:** learn homogeneous weak learners **independently** from each other in **parallel** and combine them following some kind of deterministic averaging process.
- **Boosting:** learn homogeneous weak learners **sequentially** in an adaptative way (a base model depends on the previous ones) and combine them following a deterministic strategy.
- **Stacking:** learn heterogeneous weak learners in **parallel** and combines them by **training a meta-model** to output a prediction based on the different weak models predictions.

Ref: <https://towardsdatascience.com/ensemble-methods-bagging-boosting-and-stacking-c9214a10a205>

Bagging

- Aims at producing an ensemble model that is **more robust** than the individual models composing it.
- The low correlation between models is the key.
- Bootstrapping



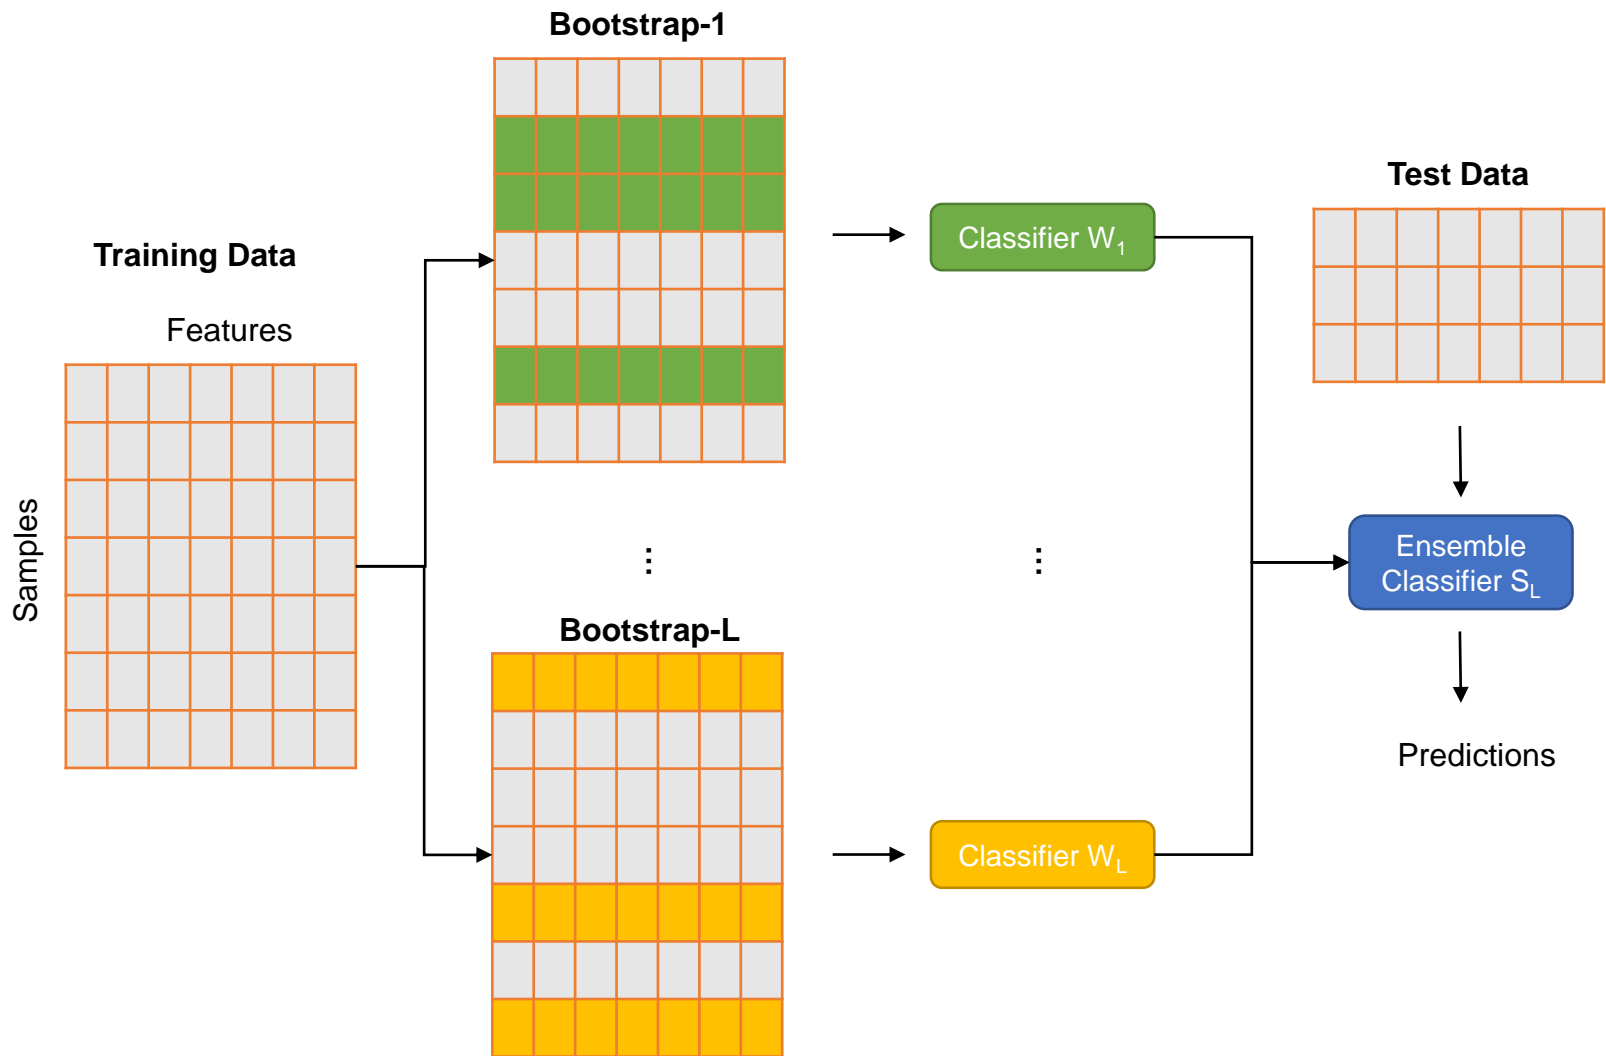
Bagging

- Assuming that we have L bootstrap samples (approximations of L independent datasets)
- We can fit L almost independent weak learners (one on each dataset)

$$w_1(.), w_2(.), \dots, w_L(.)$$

- Then aggregate them into some kind of averaging or voting process in order to get an ensemble model with a lower variance.

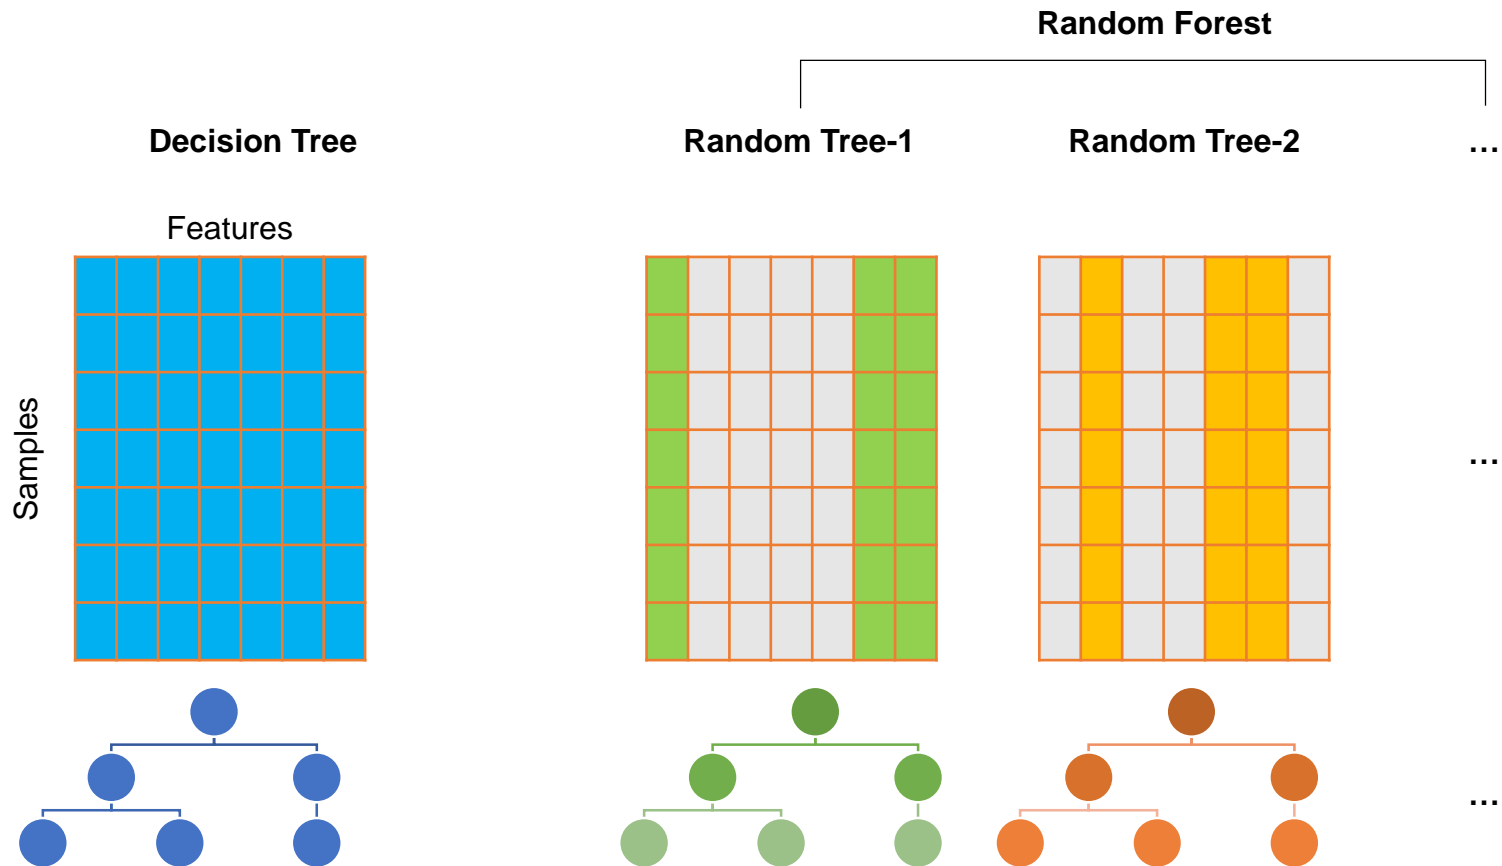
$$s_L(.) = \frac{1}{L} \sum_{l=1}^L w_l(.)$$



Random forests

- When growing each tree, instead of only sampling over the observations in the dataset to generate a bootstrap sample, it also **sample over features** and keep only a random subset of them to build the tree.
 - It forces even more variation amongst the trees in the model.
 - It reduces the correlation between the different returned outputs.
 - It makes the decision making process more robust to missing data.

Random forests



Boosting


- Main idea:
 - Train classifiers (e.g. decision trees) in a sequence.
 - A new classifier should focus on those cases which were incorrectly classified in the last round.
 - Combine the classifiers by letting them vote on the final prediction (like bagging).
 - Each classifier could be (should be) very “weak”, e.g. a decision stump.

Boosting intuition

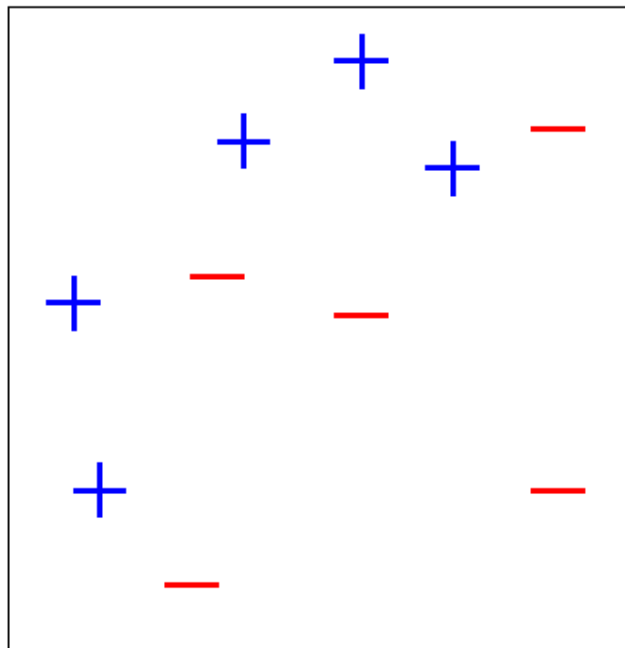
- We adaptively weigh each data case.
- Data cases which are wrongly classified get high weight (the algorithm will focus on them)
- Each boosting round learns a new (simple) classifier on the weighed dataset.
- We weigh classifiers to combine them into a single powerful classifier.
- Classifiers that obtain low training error rate have high weight (c_l).

$$\begin{array}{ccc} \text{Bagging} & & \text{Boosting} \\ s_L(.) = \sum_{l=1}^L c_l \times w_l(.) & \longrightarrow & s_l(.) = s_{l-1}(.) + c_l \times w_l(.) \end{array}$$

Boosting in a picture

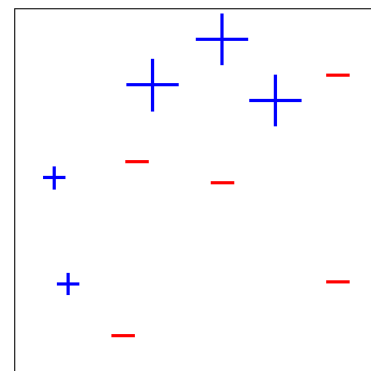
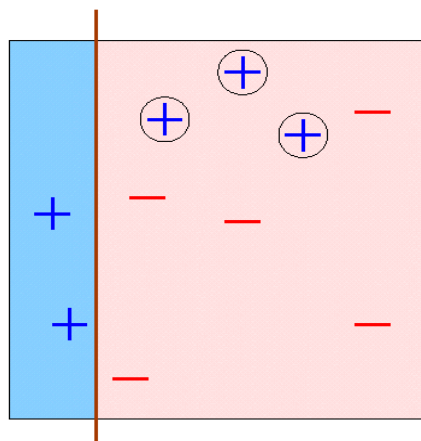


AdaBoost (Example)





Original Training set : Equal Weights to all training samples


ROUND 1




ROUND 2



ROUND 3



FINAL



AdaBoost (Algorithm)

Given: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize $D_1(i) = \frac{1}{m}$

For $t = 1, \dots, T$:

1. Find the classifier $h_t : X \rightarrow \{-1, +1\}$ that minimizes the error with respect to the distribution D_t :

$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y(i) \neq h_j(x_i)]$$

2. Prerequisite: $\epsilon_t < 0.5$, otherwise stop.

3. Choose $\alpha_t \in \mathbb{R}$, typically $\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$ where ϵ_t is the weighted error rate of classifier h_t

4. Update: $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$ where Z_t is a normalization factor

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y(i) = h_t(x_i) \\ > 1, & y(i) \neq h_t(x_i) \end{cases}$$

Output the final classifier: $H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$

AdaBoost (Algorithm)

Given: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize $D_1(i) = \frac{1}{m}$

For $t = 1, \dots, T$:

1. Find the classifier $h_t : X \rightarrow \{-1, +1\}$ that minimizes the error with respect to the distribution D_t :

$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y(i) \neq h_j(x_i)]$$

2. Prerequisite: $\epsilon_t < 0.5$, otherwise stop.

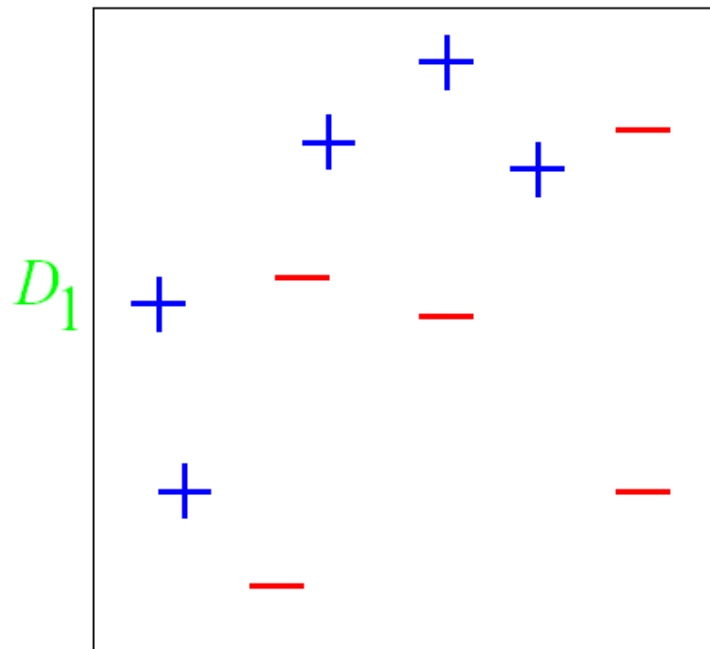
3. Choose $\alpha_t \in \mathbb{R}$, typically $\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$ where ϵ_t is the weighted error rate of classifier h_t

4. Update: $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$ where Z_t is a normalization factor

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y(i) = h_t(x_i) \\ > 1, & y(i) \neq h_t(x_i) \end{cases}$$

Output the final classifier: $H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$

AdaBoost (Example)



Original Training set : Equal Weights to all training samples

AdaBoost (Algorithm)

Given: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize $D_1(i) = \frac{1}{m}$

For $t = 1, \dots, T$:

1. Find the classifier $h_t : X \rightarrow \{-1, +1\}$ that minimizes the error with respect to the distribution D_t :

$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y(i) \neq h_j(x_i)]$$

2. Prerequisite: $\epsilon_t < 0.5$, otherwise stop.

3. Choose $\alpha_t \in \mathbb{R}$, typically $\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$ where ϵ_t is the weighted error rate of classifier h_t

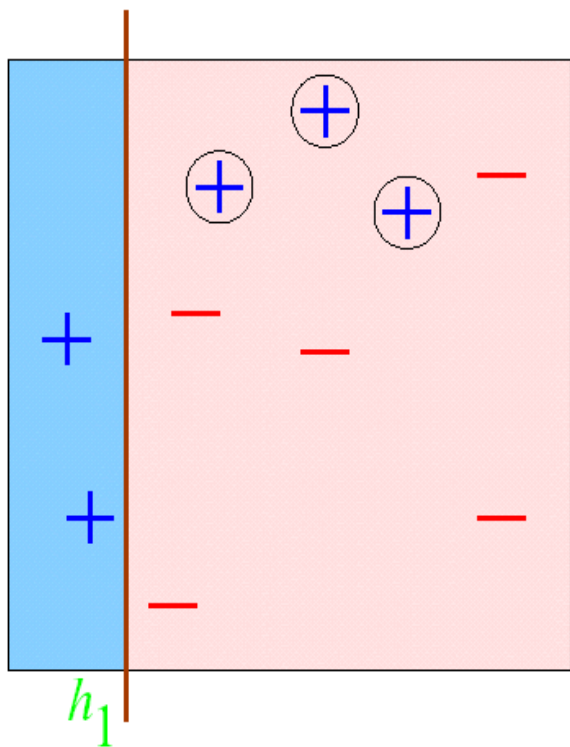
4. Update: $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$ where Z_t is a normalization factor

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y(i) = h_t(x_i) \\ > 1, & y(i) \neq h_t(x_i) \end{cases}$$

Output the final classifier: $H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$

AdaBoost (Example)


ROUND 1



$$\epsilon_1 = 0.30$$
$$\alpha_1 = 0.42$$




D_2



$$\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$$


AdaBoost (Example)

ROUND 2




$$\epsilon_2 = 0.21$$
$$\alpha_2 = 0.65$$

D_3



AdaBoost (Example)

ROUND 3



AdaBoost (Algorithm)

Given: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize $D_1(i) = \frac{1}{m}$

For $t = 1, \dots, T$:

1. Find the classifier $h_t : X \rightarrow \{-1, +1\}$ that minimizes the error with respect to the distribution D_t :

$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y(i) \neq h_j(x_i)]$$

2. Prerequisite: $\epsilon_t < 0.5$, otherwise stop.

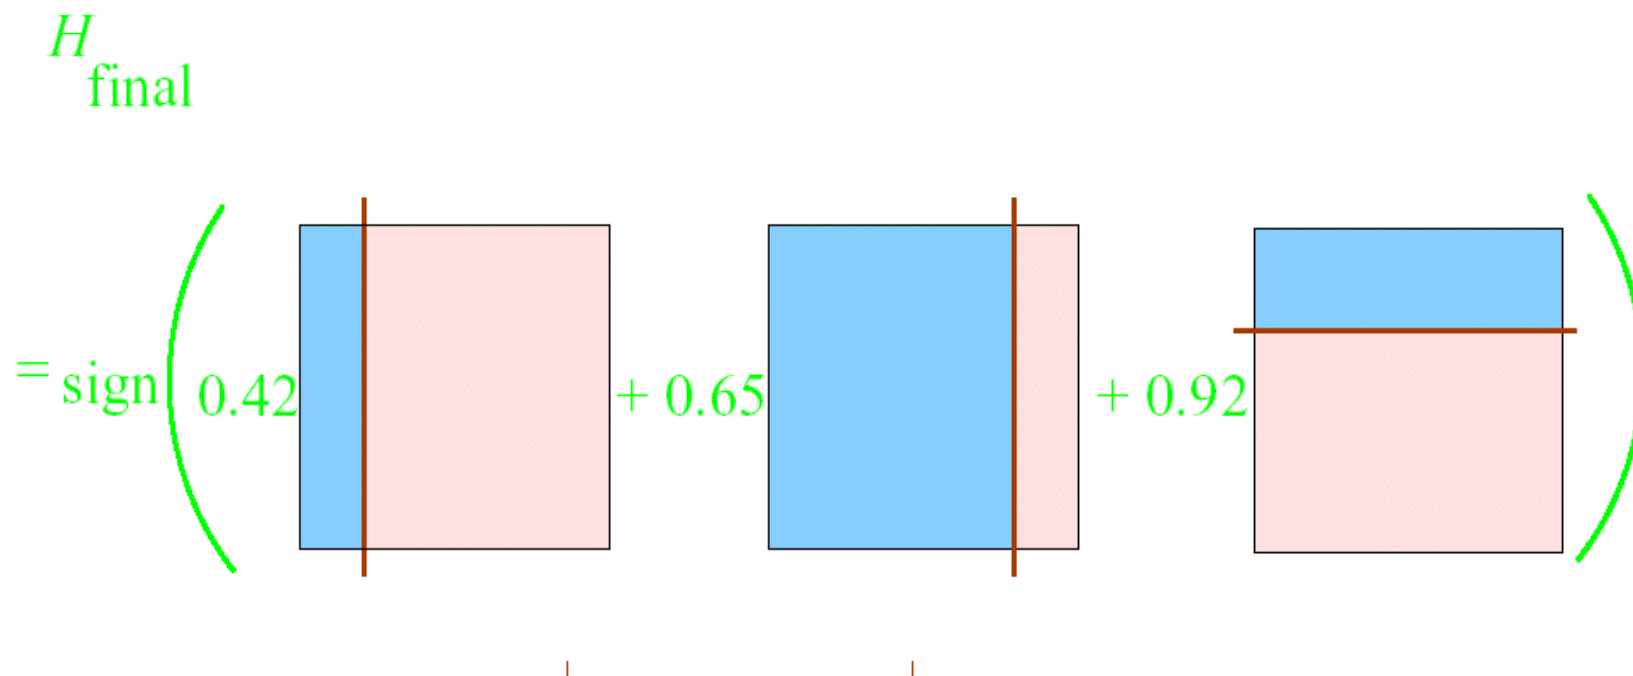
3. Choose $\alpha_t \in \mathbf{R}$, typically $\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$ where ϵ_t is the weighted error rate of classifier h_t

4. Update: $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$ where Z_t is a normalization factor

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y(i) = h_t(x_i) \\ > 1, & y(i) \neq h_t(x_i) \end{cases}$$

Output the final classifier: $H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$

AdaBoost (Example)

$$H_{\text{final}} = \text{sign} \left(0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} \right)$$


Gradient boosting

- Gradient boosting **casts the problem into a gradient descent one**: at each iteration we fit a weak learner to the opposite of the gradient of the current fitting error with respect to the current ensemble model.

Bagging

$$s_L(.) = \sum_{l=1}^L c_l \times w_l(.)$$

→

AdaBoost

$$s_l(.) = s_{l-1}(.) + c_l \times w_l(.)$$

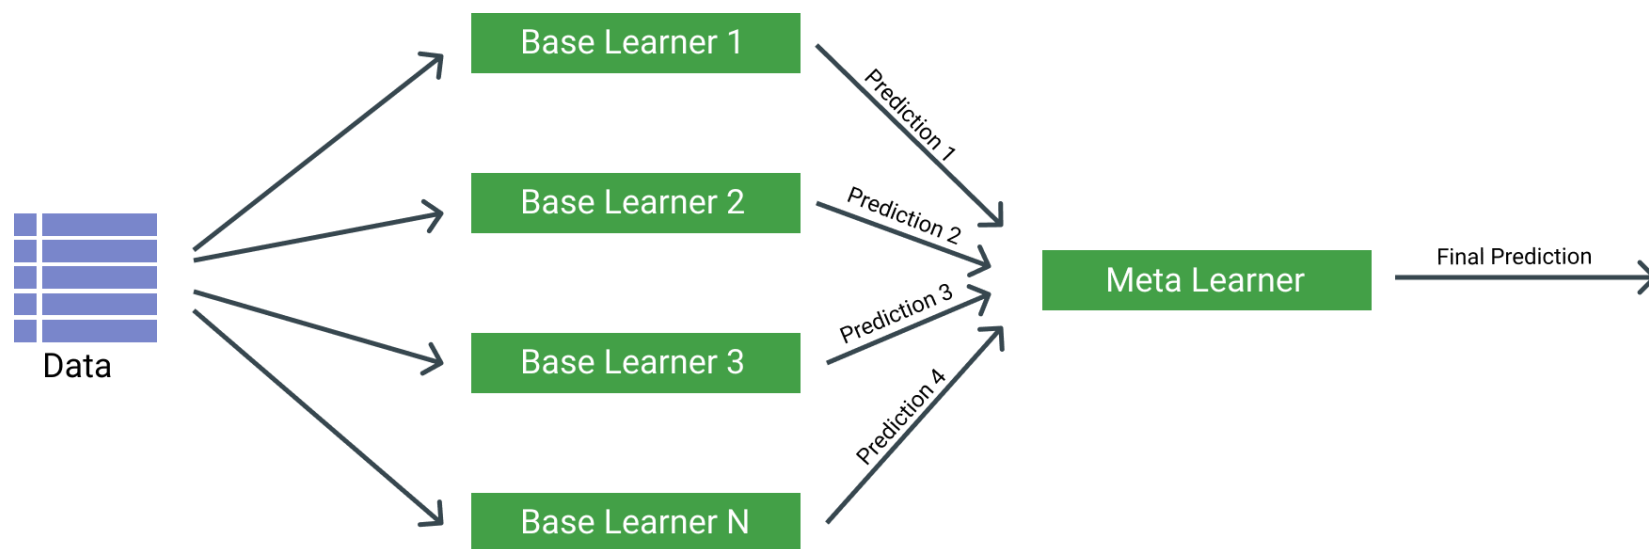
↓

$$s_l(.) = s_{l-1}(.) - c_l \times \nabla_{s_{l-1}} E(s_{l-1})(.)$$

XGBoost(Extreme Gradient Boosting)

- Additional tricks that make learning much more efficient:
 - Implements **regularization** helping reduce overfit
 - Implements **parallel** processing being much faster (10x) than GB
 - Allows users to define **custom optimization objectives** and evaluation criteria
 - XGBoost has an in-built routine to handle **missing values**
 - XGBoost **prunes** the tree backwards and removes splits beyond which there is no positive gain
 - XGBoost allows a user to run a **cross-validation at each iteration** of the boosting process

Stacking



What have we learnt

- **Decision tree**
 - Information gain
- Ensemble methods
 - Bagging
 - **Random forest**
 - Boosting
 - XGBoost
 - Stacking

Questions?



model trained
for 1000 epochs



model trained
for 100 epochs



model trained
for 1 epoch

<https://colab.research.google.com/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/05.08-Random-Forests.ipynb>